

p8106_hw1

Hao Zheng

2/20/2022

```
# Data import
train = read.csv("./data/housing_training.csv") %>% janitor::clean_names()
test = read.csv("./data/housing_test.csv") %>% janitor::clean_names()

train = na.omit(train)
test = na.omit(test)

x_train = model.matrix(sale_price~., train)[,-1]
y_train = train$sale_price

x_test <- model.matrix(sale_price~., test)[ , -1]
y_test <- test$sale_price
```

Now, let's fit different models based on the dataset.

Linear model

```
set.seed(2022)

lm.fit = lm(sale_price ~ .,
            data = train,
            method = "lm",
            trControl = trainControl(method = "repeatedcv", number = 10))
```

```
## Warning in lm(sale_price ~ ., data = train, method = "lm", trControl =
## trainControl(method = "repeatedcv", : method = 'lm' is not supported. Using 'qr'
```

```
## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'trControl' will be disregarded
```

```
summary(lm.fit)
```

```
##
## Call:
## lm(formula = sale_price ~ ., data = train, method = "lm", trControl = trainControl(method = "repeatedcv",
## number = 10))
##
## Residuals:
```

```

##      Min      1Q Median      3Q      Max
## -89864 -12424    416  12143 140205
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -4.985e+06  3.035e+06  -1.642  0.10076
## gr_liv_area      2.458e+01  1.393e+01   1.765  0.07778 .
## first_flr_sf     4.252e+01  1.409e+01   3.017  0.00260 **
## second_flr_sf    4.177e+01  1.379e+01   3.029  0.00250 **
## total_bsmt_sf    3.519e+01  2.744e+00  12.827 < 2e-16 ***
## low_qual_fin_sf          NA          NA      NA      NA
## wood_deck_sf      1.202e+01  4.861e+00   2.474  0.01350 *
## open_porch_sf     1.618e+01  1.004e+01   1.611  0.10736
## bsmt_unf_sf      -2.087e+01  1.723e+00 -12.116 < 2e-16 ***
## mas_vnr_area      1.046e+01  4.229e+00   2.473  0.01353 *
## garage_cars       4.229e+03  1.893e+03   2.234  0.02563 *
## garage_area       7.769e+00  6.497e+00   1.196  0.23195
## year_built        3.251e+02  3.130e+01  10.388 < 2e-16 ***
## tot_rms_abv_grd   -3.838e+03  6.922e+02  -5.545 3.51e-08 ***
## full_bath         -4.341e+03  1.655e+03  -2.622  0.00883 **
## overall_qualAverage -5.013e+03  1.735e+03  -2.890  0.00391 **
## overall_qualBelow_Average -1.280e+04  2.677e+03  -4.782 1.92e-06 ***
## overall_qualExcellent 7.261e+04  5.381e+03  13.494 < 2e-16 ***
## overall_qualFair    -1.115e+04  5.240e+03  -2.127  0.03356 *
## overall_qualGood     1.226e+04  1.950e+03   6.287 4.30e-10 ***
## overall_qualVery_Excellent 1.304e+05  8.803e+03  14.810 < 2e-16 ***
## overall_qualVery_Good 3.798e+04  2.741e+03  13.852 < 2e-16 ***
## kitchen_qualFair    -2.663e+04  6.325e+03  -4.210 2.71e-05 ***
## kitchen_qualGood    -1.879e+04  4.100e+03  -4.582 5.01e-06 ***
## kitchen_qualTypical -2.677e+04  4.281e+03  -6.252 5.37e-10 ***
## fireplaces          1.138e+04  2.257e+03   5.043 5.18e-07 ***
## fireplace_quFair    -7.207e+03  6.823e+03  -1.056  0.29106
## fireplace_quGood     6.070e+02  5.833e+03   0.104  0.91713
## fireplace_quNo_Fireplace 3.394e+03  6.298e+03   0.539  0.59002
## fireplace_quPoor    -5.185e+03  7.399e+03  -0.701  0.48362
## fireplace_quTypical -6.398e+03  5.897e+03  -1.085  0.27814
## exter_qualFair     -3.854e+04  8.383e+03  -4.598 4.66e-06 ***
## exter_qualGood     -1.994e+04  5.585e+03  -3.569  0.00037 ***
## exter_qualTypical   -2.436e+04  5.874e+03  -4.147 3.57e-05 ***
## lot_frontage        1.024e+02  1.905e+01   5.376 8.90e-08 ***
## lot_area            6.042e-01  7.864e-02   7.683 2.91e-14 ***
## longitude          -3.481e+04  2.537e+04  -1.372  0.17016
## latitude            5.874e+04  3.483e+04   1.686  0.09193 .
## misc_val            9.171e-01  1.003e+00   0.914  0.36071
## year_sold          -6.455e+02  4.606e+02  -1.401  0.16132
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 22190 on 1401 degrees of freedom
## Multiple R-squared:  0.9116, Adjusted R-squared:  0.9092
## F-statistic: 380.3 on 38 and 1401 DF, p-value: < 2.2e-16

```

```
pred.lm <- predict(lm.fit, newdata = test)
```

```
## Warning in predict.lm(lm.fit, newdata = test): prediction from a rank-deficient
## fit may be misleading
```

```
lm_mse = RMSE(pred.lm, test$sale_price); lm_mse
```

```
## [1] 21149.18
```

The test MSE for the least square method is 2.1149176×10^4 . Potential disadvantage of the linear model:
1. There may be too many predictors, which could cause problems such as corlinearity among predictors, large variance; 2. The model is too complex and there exist over-fitting problems.


Lasso model

using glmnet


```
set.seed(2022)

lasso.fit <- cv.glmnet(
  x = x_train,
  y = y_train,
  alpha = 1,
  lambda = exp(seq(8, 3, length = 100))
)

plot(lasso.fit)
```



```
plot_glmnet(lasso.fit$glmnet.fit)
```



```
# Look at the 1SE coefficient for lasso
predict(lasso.fit, s = "lambda.1se", type = "coefficients")
```

```
## 40 x 1 sparse Matrix of class "dgCMatrix"
##               lambda.1se
## (Intercept)      -3.411186e+06
## gr_liv_area       5.916147e+01
## first_flr_sf      9.908879e-01
## second_flr_sf      .
## total_bsmt_sf     3.660286e+01
## low_qual_fin_sf   -3.110661e+01
## wood_deck_sf      9.211416e+00
## open_porch_sf     1.040501e+01
## bsmt_unf_sf       -2.036685e+01
## mas_vnr_area      1.368374e+01
## garage_cars       3.296289e+03
## garage_area       1.014693e+01
## year_built        3.116010e+02
## tot_rms_abv_grd   -2.051953e+03
## full_bath         -3.535970e+02
## overall_qualAverage -3.577405e+03
## overall_qualBelow_Average -1.014004e+04
## overall_qualExcellent 9.096805e+04
## overall_qualFair   -7.842522e+03
## overall_qualGood    1.067898e+04
## overall_qualVery_Excellent 1.608334e+05
```

```
## overall_qualVery_Good      3.716056e+04
## kitchen_qualFair          -7.690510e+03
## kitchen_qualGood          -1.468331e+03
## kitchen_qualTypical       -1.076409e+04
## fireplaces                 7.195489e+03
## fireplace_quFair          -1.473565e+03
## fireplace_quGood           3.625170e+03
## fireplace_quNo_Fireplace   .
## fireplace_quPoor           .
## fireplace_quTypical        -2.188237e+03
## exter_qualFair            -1.590797e+04
## exter_qualGood             .
## exter_qualTypical          -4.741677e+03
## lot_frontage               8.043093e+01
## lot_area                   5.822890e-01
## longitude                  -1.729013e+04
## latitude                   2.980870e+04
## misc_val                   2.214766e-02
## year_sold                  -1.300149e+01
```

When the 1SE rule is applied, we can see there are 35 predictors included in the model.

```
y_pred <- predict(lasso.fit, newx = x_test, s = "lambda.min", type = "response")
lasso_mse <- mean(RMSE(y_pred, y_test)^2); lasso_mse
```

```
## [1] 440064267
```

The test MSE for Lasso model (1SE) is 4.4006427×10^8 .

Elastic Net model

```
set.seed(2022)

enet.fit = train(x_train, y_train,
                 method = "glmnet",
                 tuneGrid = expand.grid(alpha = seq(0, 1, length = 21),
                                       lambda = exp(seq(8, -3, length = 50))),
                 trControl = trainControl(method = "cv", number = 10))


enet.fit$bestTune

##      alpha      lambda
## 95  0.05 970.2473
```

The selected tuning parameter $\lambda = 970.247332$, $\alpha = 0.05$. Then we visualize the elastic net result.

```
# Visualization
myCol <- rainbow(25)
myPar <- list(superpose.symbol = list(col = myCol),
              superpose.line = list(col = myCol))

plot(enet.fit, par.settings = myPar)
```



```
coef(enet.fit$finalModel, enet.fit$bestTune$lambda)
```

```
## 40 x 1 sparse Matrix of class "dgCMatrix"
##               s1
## (Intercept)   -5.172951e+06
## gr_liv_area    3.812434e+01
## first_flr_sf   2.637357e+01
## second_flr_sf  2.480921e+01
## total_bsmt_sf  3.473579e+01
## low_qual_fin_sf -1.604315e+01
## wood_deck_sf   1.252592e+01
## open_porch_sf  1.727662e+01
## bsmt_unf_sf    -2.057836e+01
## mas_vnr_area   1.235856e+01
## garage_cars    3.958861e+03
## garage_area    9.528576e+00
## year_built     3.152748e+02
## tot_rms_abv_grd -3.197583e+03
## full_bath      -3.313259e+03
## overall_qualAverage -5.160203e+03
## overall_qualBelow_Average -1.263204e+04
## overall_qualExcellent 7.695123e+04
## overall_qualFair -1.165139e+04
## overall_qualGood 1.174884e+04
## overall_qualVery_Excellent 1.386447e+05
```

```
## overall_qualVery_Good      3.733755e+04
## kitchen_qualFair          -2.220863e+04
## kitchen_qualGood          -1.475407e+04
## kitchen_qualTypical       -2.285583e+04
## fireplaces                 1.053606e+04
## fireplace_quFair          -8.000449e+03
## fireplace_quGood           6.186069e+01
## fireplace_quNo_Fireplace   1.117703e+03
## fireplace_quPoor          -5.936463e+03
## fireplace_quTypical       -7.046712e+03
## exter_qualFair            -3.077995e+04
## exter_qualGood            -1.239462e+04
## exter_qualTypical         -1.717937e+04
## lot_frontage               9.877631e+01
## lot_area                   6.017072e-01
## longitude                  -3.534682e+04
## latitude                   5.708440e+04
## misc_val                   8.336667e-01
## year_sold                  -5.374119e+02
```

Next, we calculate the test MSE on our test dataset.

```
# Elastic net test MSE
enet_pred <- predict(enet.fit, newdata = x_test)
enet_mse <- mean(RMSE(enet_pred, y_test)^2); enet_mse
```

```
## [1] 434592441
```

The test error for Elastic net model is 4.3459244×10^8 .

Partial Least Square model

```
set.seed(2022)

pls.fit <- plsr(sale_price~.,
               data = train,
               scale = TRUE,
               validation = "CV")

summary(pls.fit)

## Data:      X dimension: 1440 39
## Y dimension: 1440 1
## Fit method: kernelpls
## Number of components considered: 39
##
## VALIDATION: RMSEP
## Cross-validated using 10 random segments.
##      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps
## CV           73685   33432   27986   25125   24011   23369   23171
```




```

## adjCV      73685    33426    27949    25054    23942    23303    23113
##           7 comps  8 comps  9 comps 10 comps 11 comps 12 comps 13 comps
## CV        23078    23036    23033    23027    23027    23014    23009
## adjCV      23022    22982    22977    22969    22967    22955    22950
##           14 comps 15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
## CV        23013    23020    23021    23025    23030    23027    23032
## adjCV      22954    22960    22961    22965    22969    22966    22971
##           21 comps 22 comps 23 comps 24 comps 25 comps 26 comps 27 comps
## CV        23032    23032    23032    23032    23032    23032    23034
## adjCV      22971    22971    22971    22971    22971    22971    22973
##           28 comps 29 comps 30 comps 31 comps 32 comps 33 comps 34 comps
## CV        23034    23034    23034    23034    23034    23034    23034
## adjCV      22973    22973    22973    22973    22973    22973    22973
##           35 comps 36 comps 37 comps 38 comps 39 comps
## CV        23034    23034    23034    23034    24005
## adjCV      22973    22973    22973    22973    23330
##
## TRAINING: % variance explained
##           1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps
## X           20.02    25.93    29.67    33.59    37.01    40.03    42.49
## sale_price   79.73    86.35    89.36    90.37    90.87    90.99    91.06
##           8 comps  9 comps 10 comps 11 comps 12 comps 13 comps 14 comps
## X           45.53    47.97    50.15    52.01    53.69    55.35    56.86
## sale_price   91.08    91.10    91.13    91.15    91.15    91.16    91.16
##           15 comps 16 comps 17 comps 18 comps 19 comps 20 comps
## X           58.64    60.01    62.18    63.87    65.26    67.10
## sale_price   91.16    91.16    91.16    91.16    91.16    91.16
##           21 comps 22 comps 23 comps 24 comps 25 comps 26 comps
## X           68.44    70.12    71.72    73.35    75.20    77.27
## sale_price   91.16    91.16    91.16    91.16    91.16    91.16
##           27 comps 28 comps 29 comps 30 comps 31 comps 32 comps
## X           78.97    80.10    81.83    83.55    84.39    86.34
## sale_price   91.16    91.16    91.16    91.16    91.16    91.16
##           33 comps 34 comps 35 comps 36 comps 37 comps 38 comps
## X           88.63    90.79    92.79    95.45    97.49    100.00
## sale_price   91.16    91.16    91.16    91.16    91.16    91.16
##           39 comps
## X           100.67
## sale_price   91.16

```

```
validationplot(pls.fit, val.type="MSEP", legendpos = "topright")
```



```
# Calculate the number of component in the model
cv.mse <- RMSEP(pls.fit)
ncomp.cv <- which.min(cv.mse$val[1,,]) - 1; ncomp.cv
```

```
## 13 comps
##      13
```

```
pls_pred <- predict(pls.fit, newdata = x_test, ncomp = ncomp.cv)
pls_mse <- mean(RMSE(y_test, pls_pred)^2); pls_mse
```

```
## [1] 448737340
```

There are 13 component in pls model, and the test error (MSE) is 4.4873734×10^8 .

Comparing different models

```
name <- c("lm", "lasso 1se", "elastic net", "pls")
MSE <- c(lm_mse, lasso_mse,enet_mse, pls_mse)
comparison <- cbind(name, MSE)
comparison <- as.data.frame(comparison)
```

Now, let's compare the test MSE of the above 4 models, as we have mentioned before, linear model may have many disadvantages, so here though the test MSE for linear model is the lowest, we may tend to choose

other models. Therefore, we may choose the model with the lowest MSE besides the linear model: elastic net model.